

1.

(a).

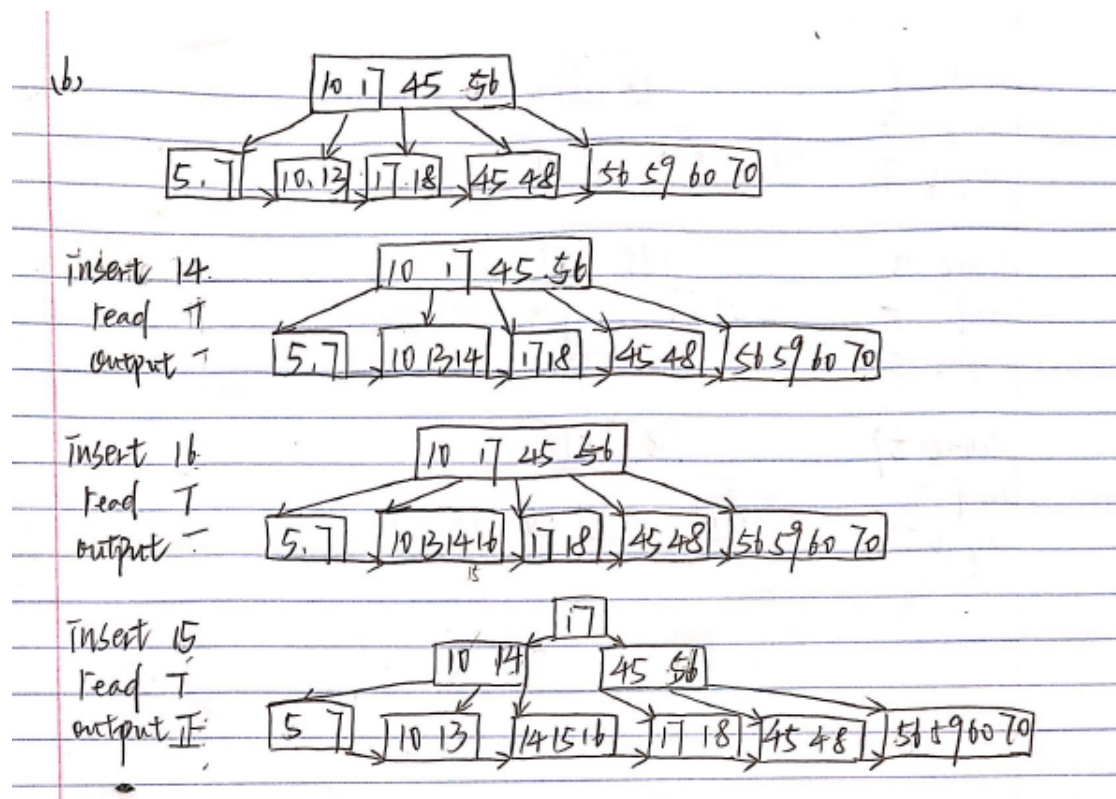
First, going to the root. --- 1 block

Second, going to the second child where value of node is beginning from 10 --- 1 block

Third, searching till the value of node is 48, --- 2 blocks.

Then, read the 5th child to estimate which is >50, --- 1 block


**Total cost = 5 blocks.**




(b). **Total cost = 13 blocks**

(c)


delete 5  
read IF  
output T




delete 7  
read T  
output -




delete 13  
read F  
output T




delete 15  
read T  
output -




delete 17  
read F  
output T



delete 45  
read IF  
output F



delete 59  
read T  
output -



(c). Total cost = 33 blocks

2.

(a) nested-loop join

1. For each  $(M-2)=100$  blocks in R (read cost =  $B(R)$ )

2. iterate through every block in S;
3. iterate through each tuple in R block,
4. iterate through each tuple in S block,
5. if condition matches, return (r, s).

Run this algorithm  $B(R)/(M-2)$  times, and each time need to read S: costs =  $B(S) * B(R) / (M-2)$ .

$$\text{Total cost} = B(R) + B(R)B(S)/(M-2) = 20\,000 + (20\,000 * 10\,000)/100 = 2,020,000$$

**(b)** nested-loop join

1. For each  $(M-2)=100$  blocks in S (read cost= $B(S)$ ),
2. iterate through every block in R;
3. iterate through each tuple in S block
4. iterate through each tuple in R block,
5. if condition matches, return (s, r).

Run this algorithm  $B(S)/(M-2)$  times, and each time need to read R: costs =  $B(S) * B(R) / (M-2)$ .

$$\text{Total cost} = B(S) + B(R)B(S)/(M-2) = 10\,000 + (20\,000 * 10\,000)/100 = 2,010,000$$

**(c)** sort-merge join

Step 1: Sort R: load 100 blocks of R at a time, sort them and write back on disk. This generates **200 runs**, each of size 100 blocks. Cost= $2B(R)=40000$

Step 2: Sort S: load 100 blocks of S at a time, sort them and write back on disk. This generates **100 runs**, each of size 100 blocks. Cost= $2B(S)=20,000$

Since # of runs of R and S  $> M$ , we cannot merge them directly.

Step 3: Sort R: load 100 blocks of R at a time, sort them and write back on disk. This generates **2 runs** Cost= $2B(R)=40000$

Since # of runs of R and S  $= 102 > M=101$ , we cannot merge them directly.

Step 4: Sort R: load 100 blocks of s at a time, sort them and write back on disk. This generates **1 runs** Cost= $2B(S)=20000$

Step 4: Final merge: Cost =  $B(R)+B(S) = 30\,000$

$$\text{Total cost} = 150000$$

**(d)** simple sort-merge join

Step 1: Sort R: load 100 blocks of R at a time, sort them and write back on disk. This generates **200 runs**, each of size 100 blocks. Cost=2B(R)=40000

Step 2: Sort S: load 100 blocks of S at a time, sort them and write back on disk. This generates **100 runs**, each of size 100 blocks. Cost=2B(S)=20,000

Step 3: Load 100 blocks of R at a time, sort and write back to the disk. This generates **2 runs**, each of size 10000 blocks. Cost =  $2B(R) = 40000$

Step 4: Load 100 blocks of S at a time, sort and write back to the disk. This generate **1 runs**, each of size 10000 blocks Cost=2B(S)=20,000

Step 5: Load 100 blocks of R at a time, sort and write back to the disk. This generate **1 runs**, each of size 20000 blocks. Cost =  $2B(R) = 40000$

Step 6: Final merge: Cost = B(R)+B(S) = 30 000

**Total cost = 190000**

**(e)** Partitioned-hash join

Step1: Hash R into  $M - 1 = 101$  buckets and send them back to disk: Cost =  $2 B(R) = 40,000$

Step2: Hash S into  $M - 1 = 101$  buckets and send them back to disk: Cost =  $2 B(S) = 20,000$

Step3: Read each bucket of smaller relation into memory. For each bucket, read the same bucket of the: larger relation block by block and join the matching tuples. Cost:  $B(R) + B(S) = 30,000$

**Total cost =  $3B(R) + 3B(S) = 90,000$**

**(f)** Index-based join

S, R are clustered

1. Load 100 blocks from R
2. For each tuple  $r(a.b)$ , fetch corresponding tuples from S
3. Join the tuples which are matched and send it to output buffer

**Total cost =  $B(R) + T(R)B(S)/v(s,a) = 200\ 00 + 200000*10000/100 = 20,020,000$**

**The best one is Partitioned-hash join.**